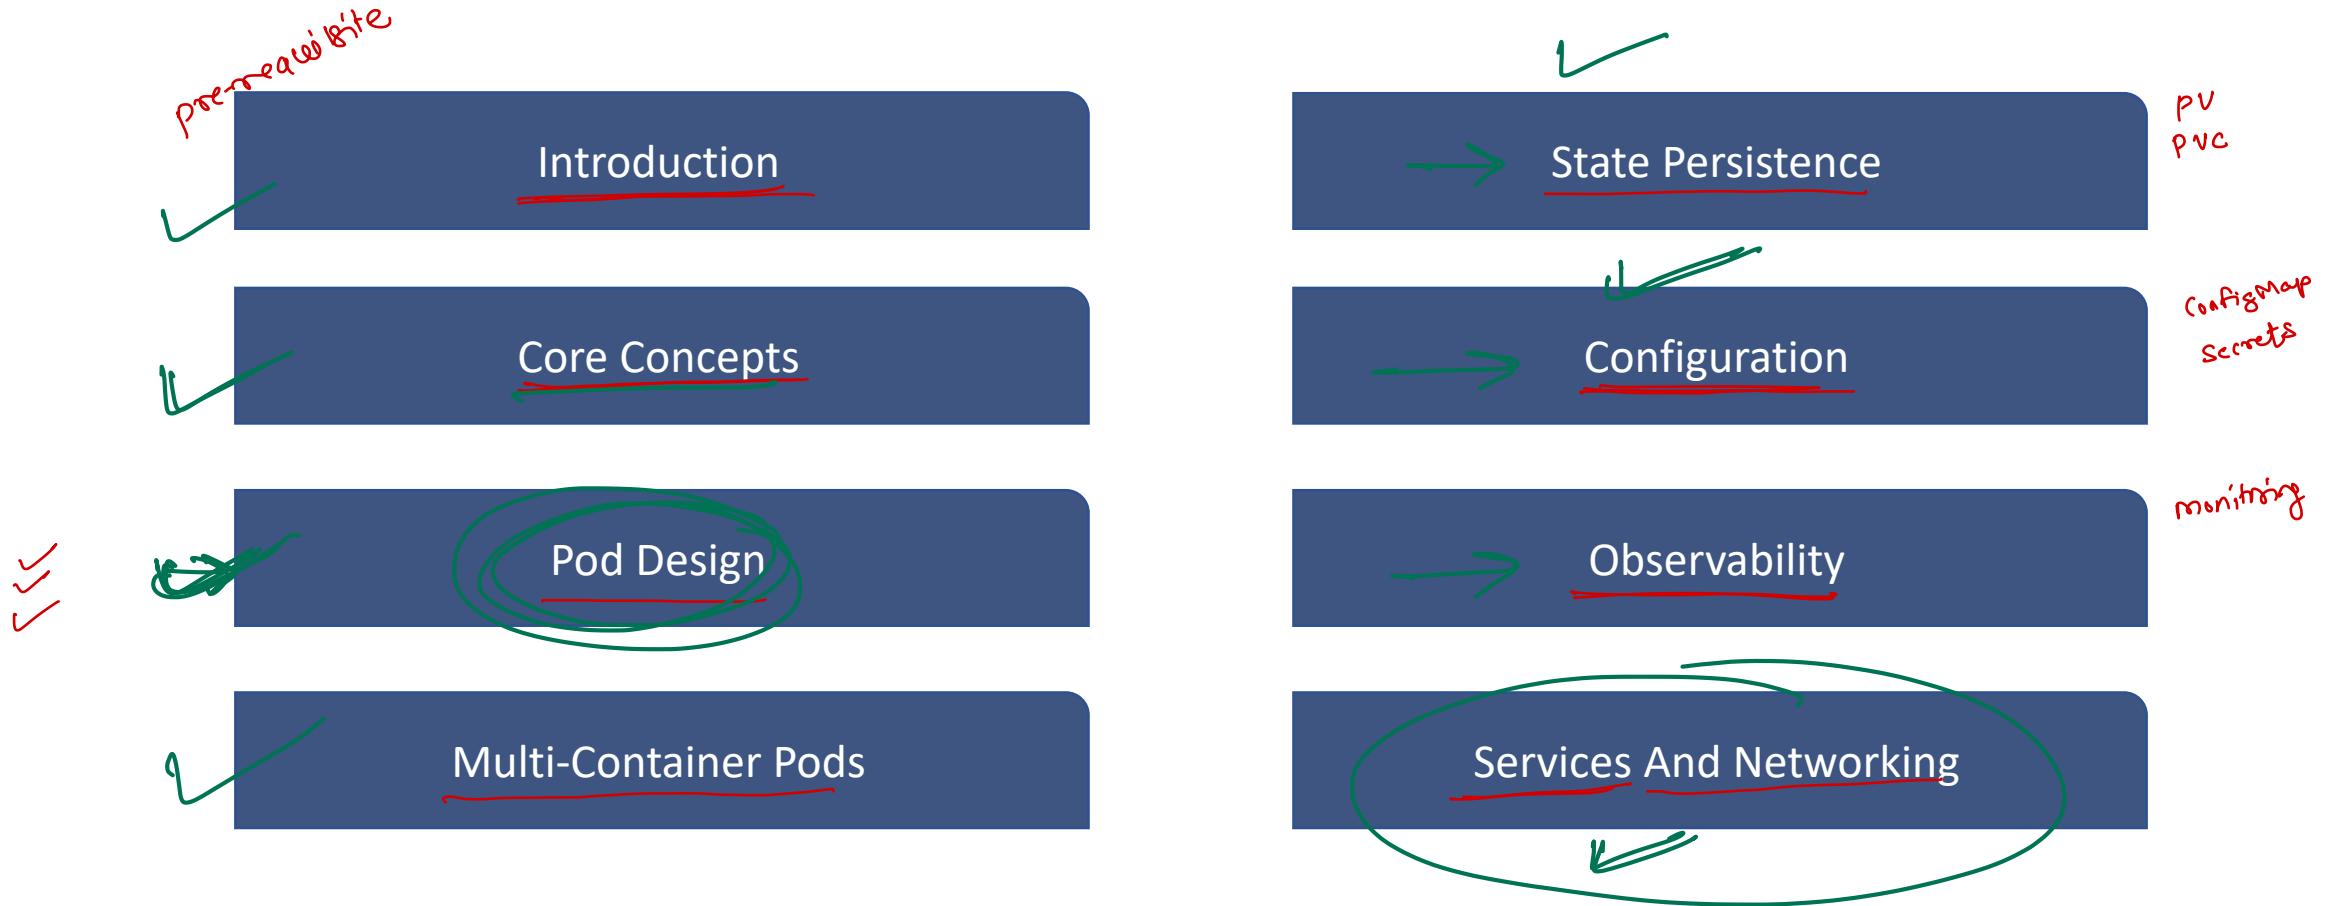




kubernetes







Cluster Installation & Configuration

①

Cluster Administration

daemonset
✓ ReplicaSet
✓ Replication Controller
✓ Deployment
✱

Workloads and Scheduling

②

Troubleshooting

Storage

PV
PVC
Storage classes
✱

Services & Networking

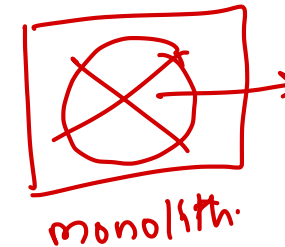


Pre-requisites




Microservice

- Distinctive method of developing software systems that tries to focus on building single-function modules with well-defined interfaces and operations
- Is an architectural style that structures an application as a collection of services that are
 - Highly maintainable and testable
 - Loosely coupled
 - Independently deployable
 - Organized around business capabilities

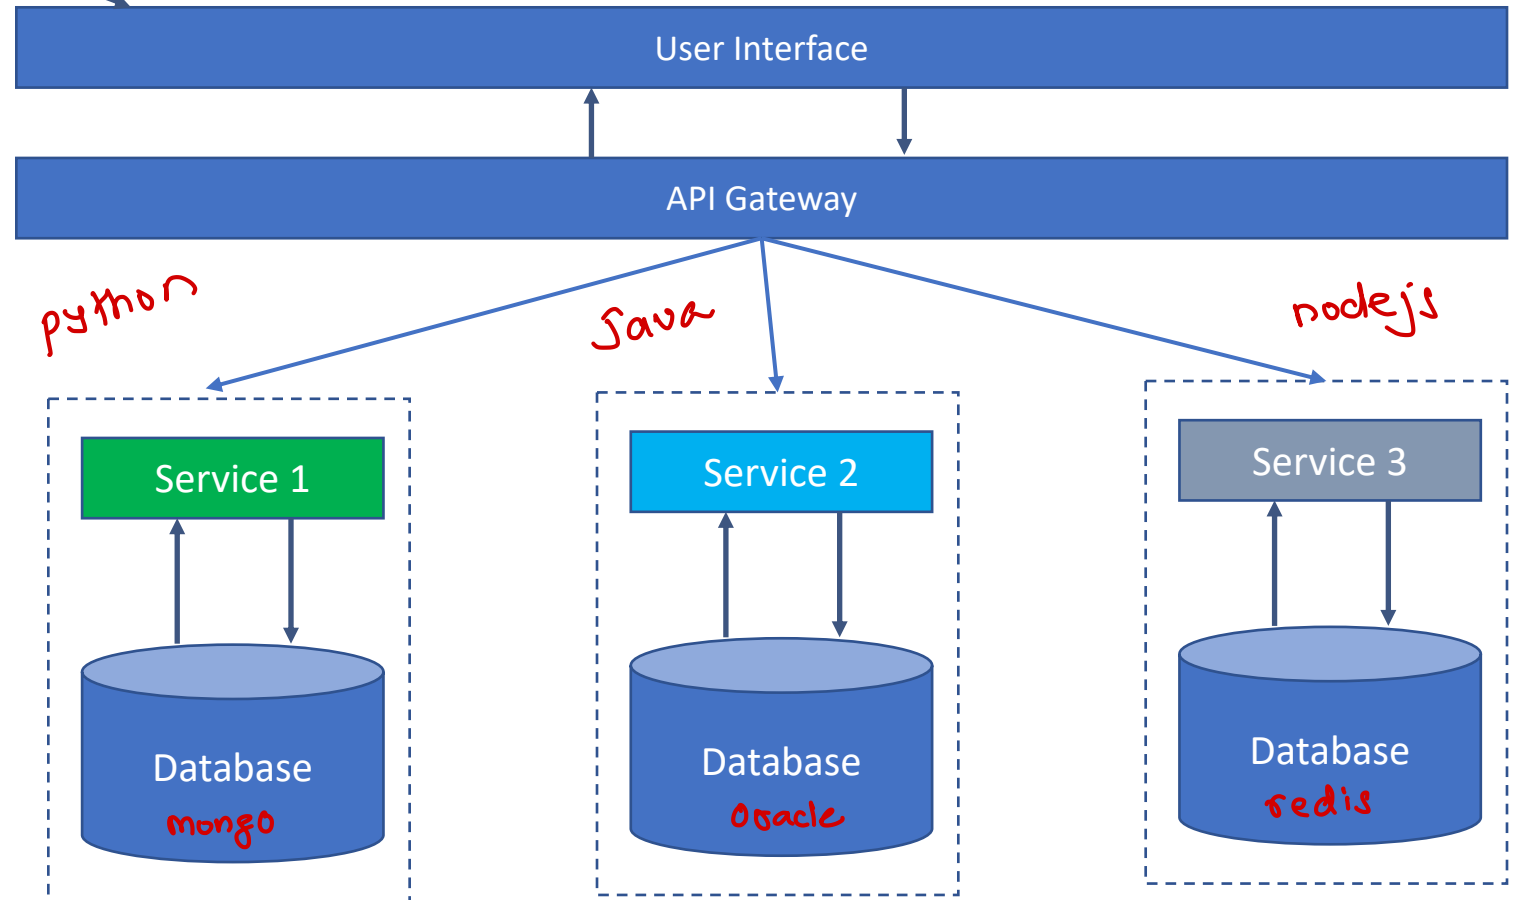


Monolithic vs Microservice


Monolithic




Microservices



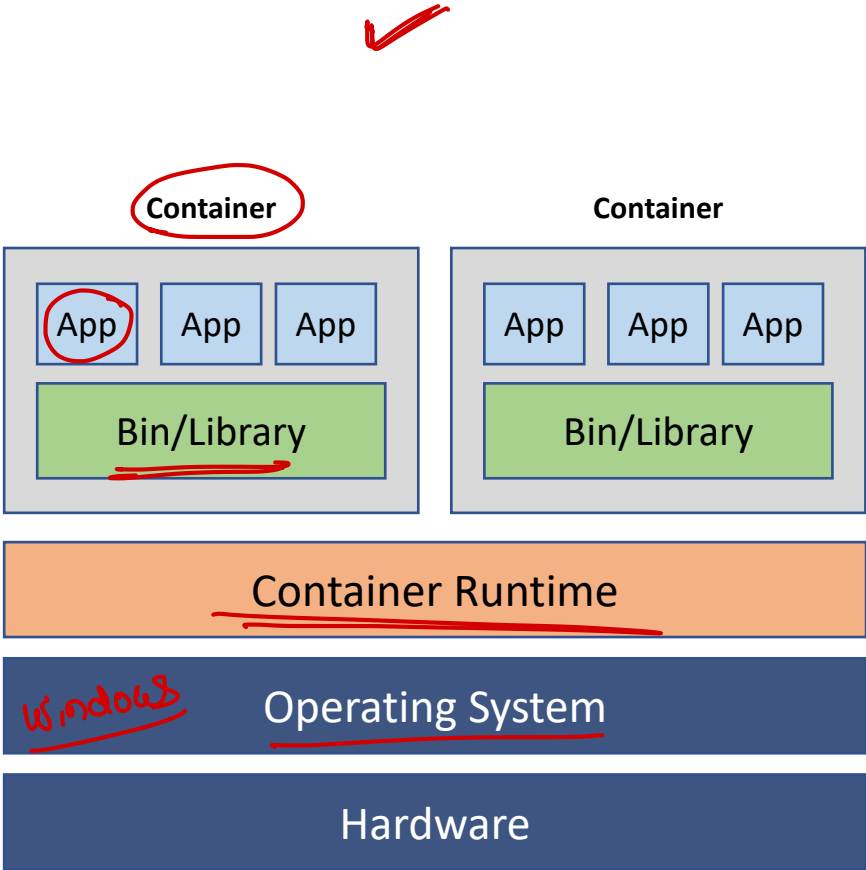
Deployment options



Traditional Deployment



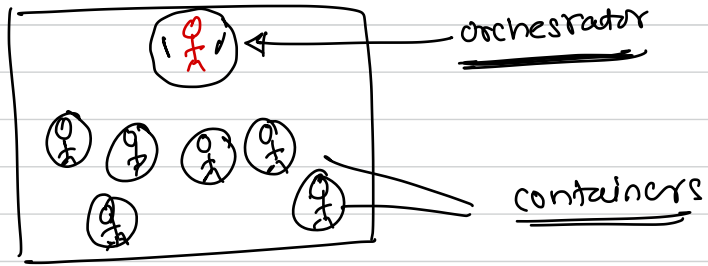
Virtual Machine Deployment



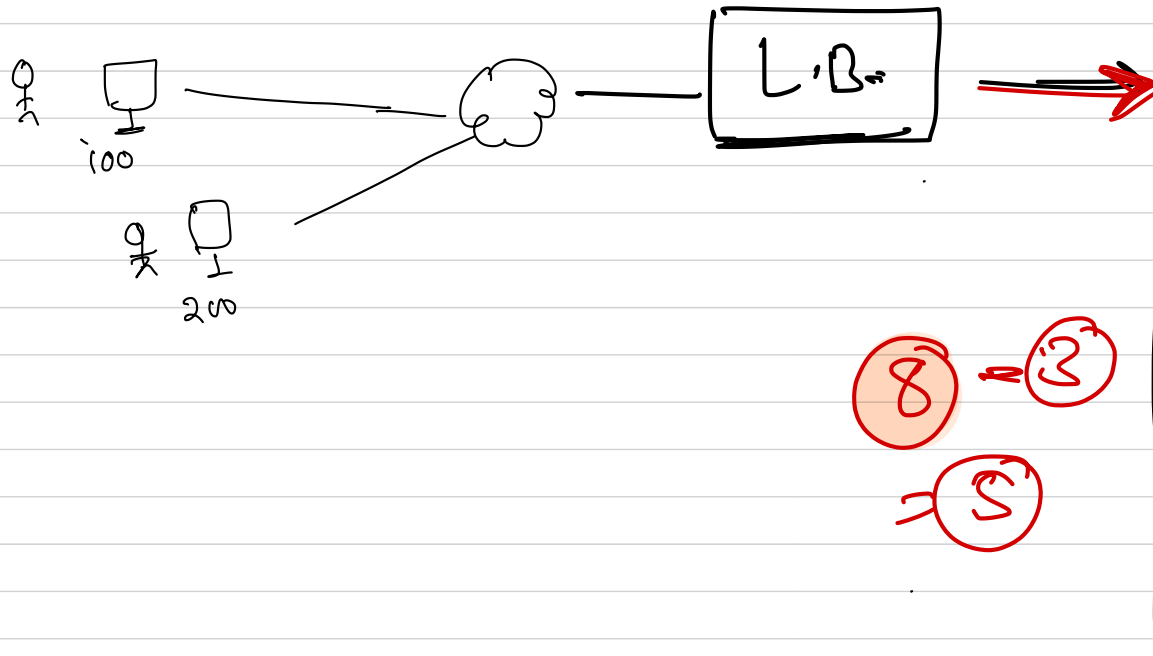
Containerized Deployment



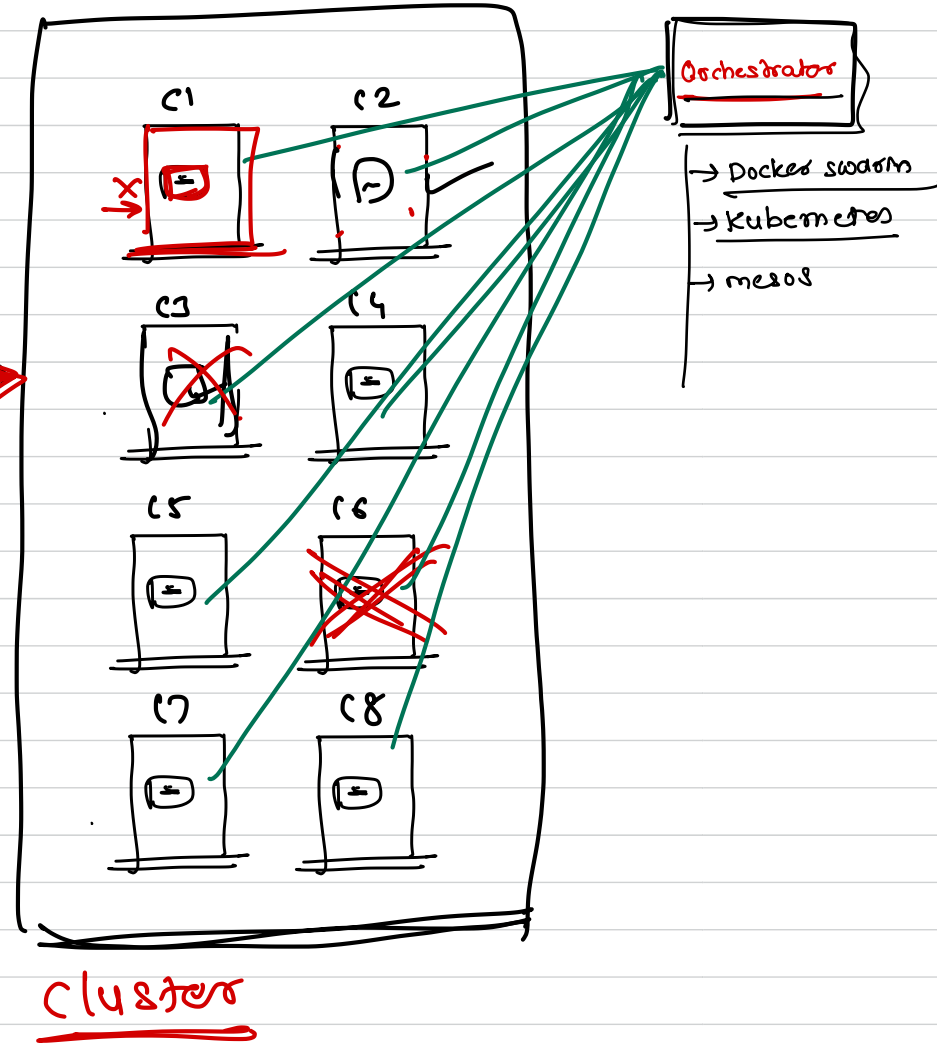
orchestra

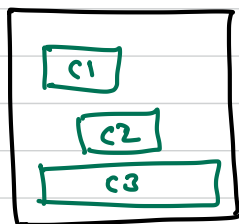


Horizontal scaling

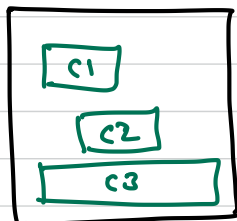


8 3 5

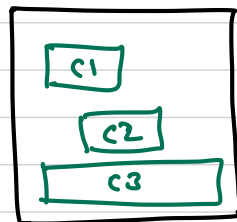





EC2 (VM)



EC2 (VM)



EC2 (VM)



EC2 (VM)

Container orchestration

- Container orchestration is all about managing the lifecycles of containers
- Software teams use container orchestration to control and automate many tasks
 - Provisioning and deployment of containers
 - Redundancy and availability of containers
 - ✓ ■ Scaling up or removing containers to spread application load evenly across host infrastructure
 - Allocation of resources between containers
 - External exposure of services running in a container with the outside world
 - Load balancing of service discovery between containers
 - Health monitoring of containers and hosts
 - Configuration of an application in relation to the containers running it
 - Movement of containers from one host to another if there is a shortage of resources, or if a host dies
- Tools
 - Docker Swarm
 - ✓ ■ Kubernetes
 - Apache Mesos and Marathon



Introduction

- Kubernetes is an open source system for automating deployments, scaling and management of containerized applications
- It is a container orchestration system
- It enables organizations to automate deployment and manage the containers, thus helping them to streamline and simplify the day to day routines
- If you are using Kubernetes that means your application is following the microservices architecture and is already containerized



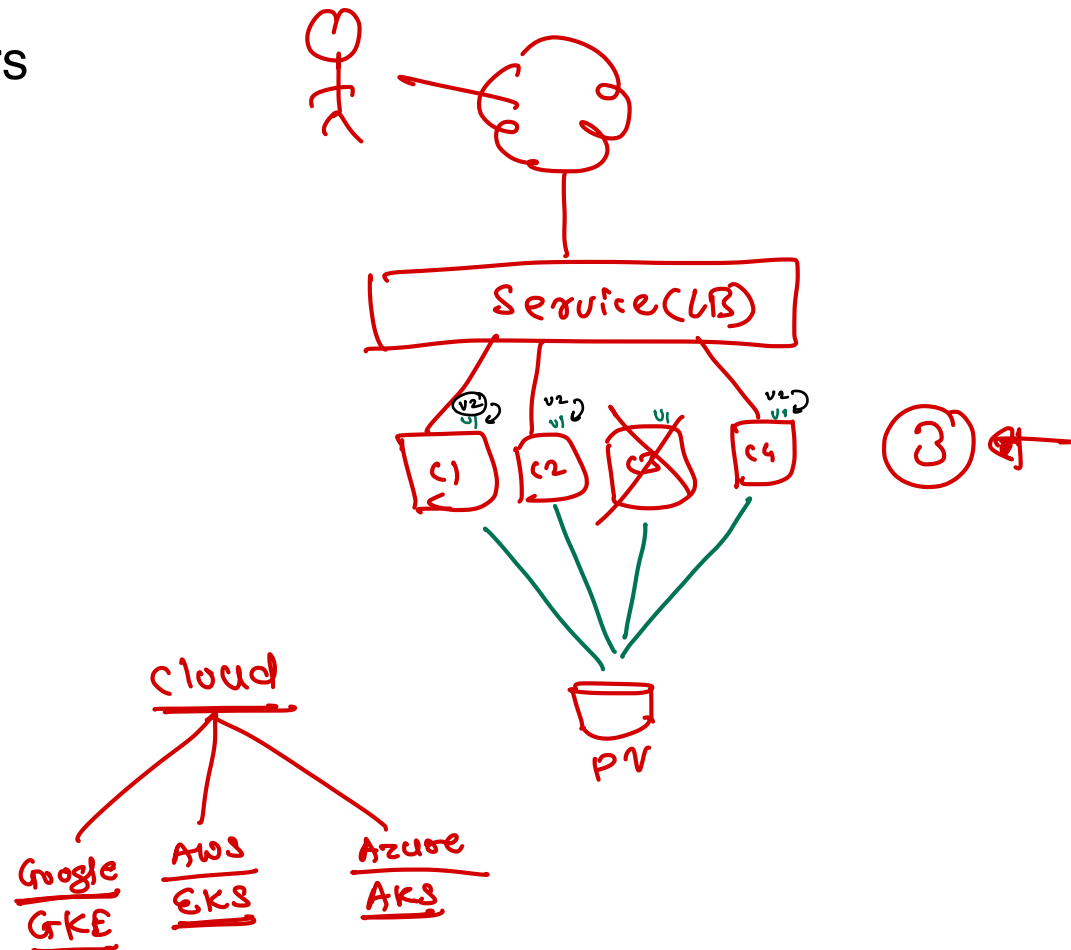
History

- Kubernetes was originally developed at Google having emerged from project Borg
- Now it has been taken over by Cloud Native Computing Foundation (CNCF)
- It is not only open source but it is managed by open community



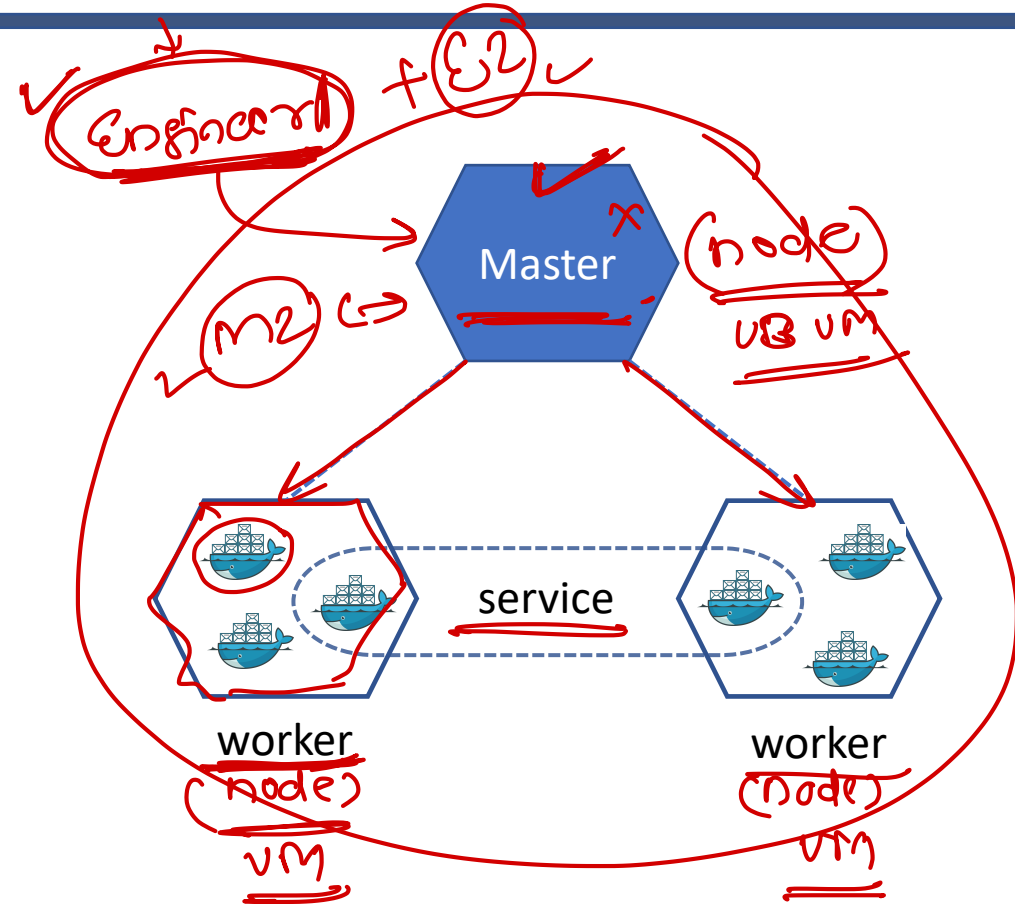
Kubernetes Features

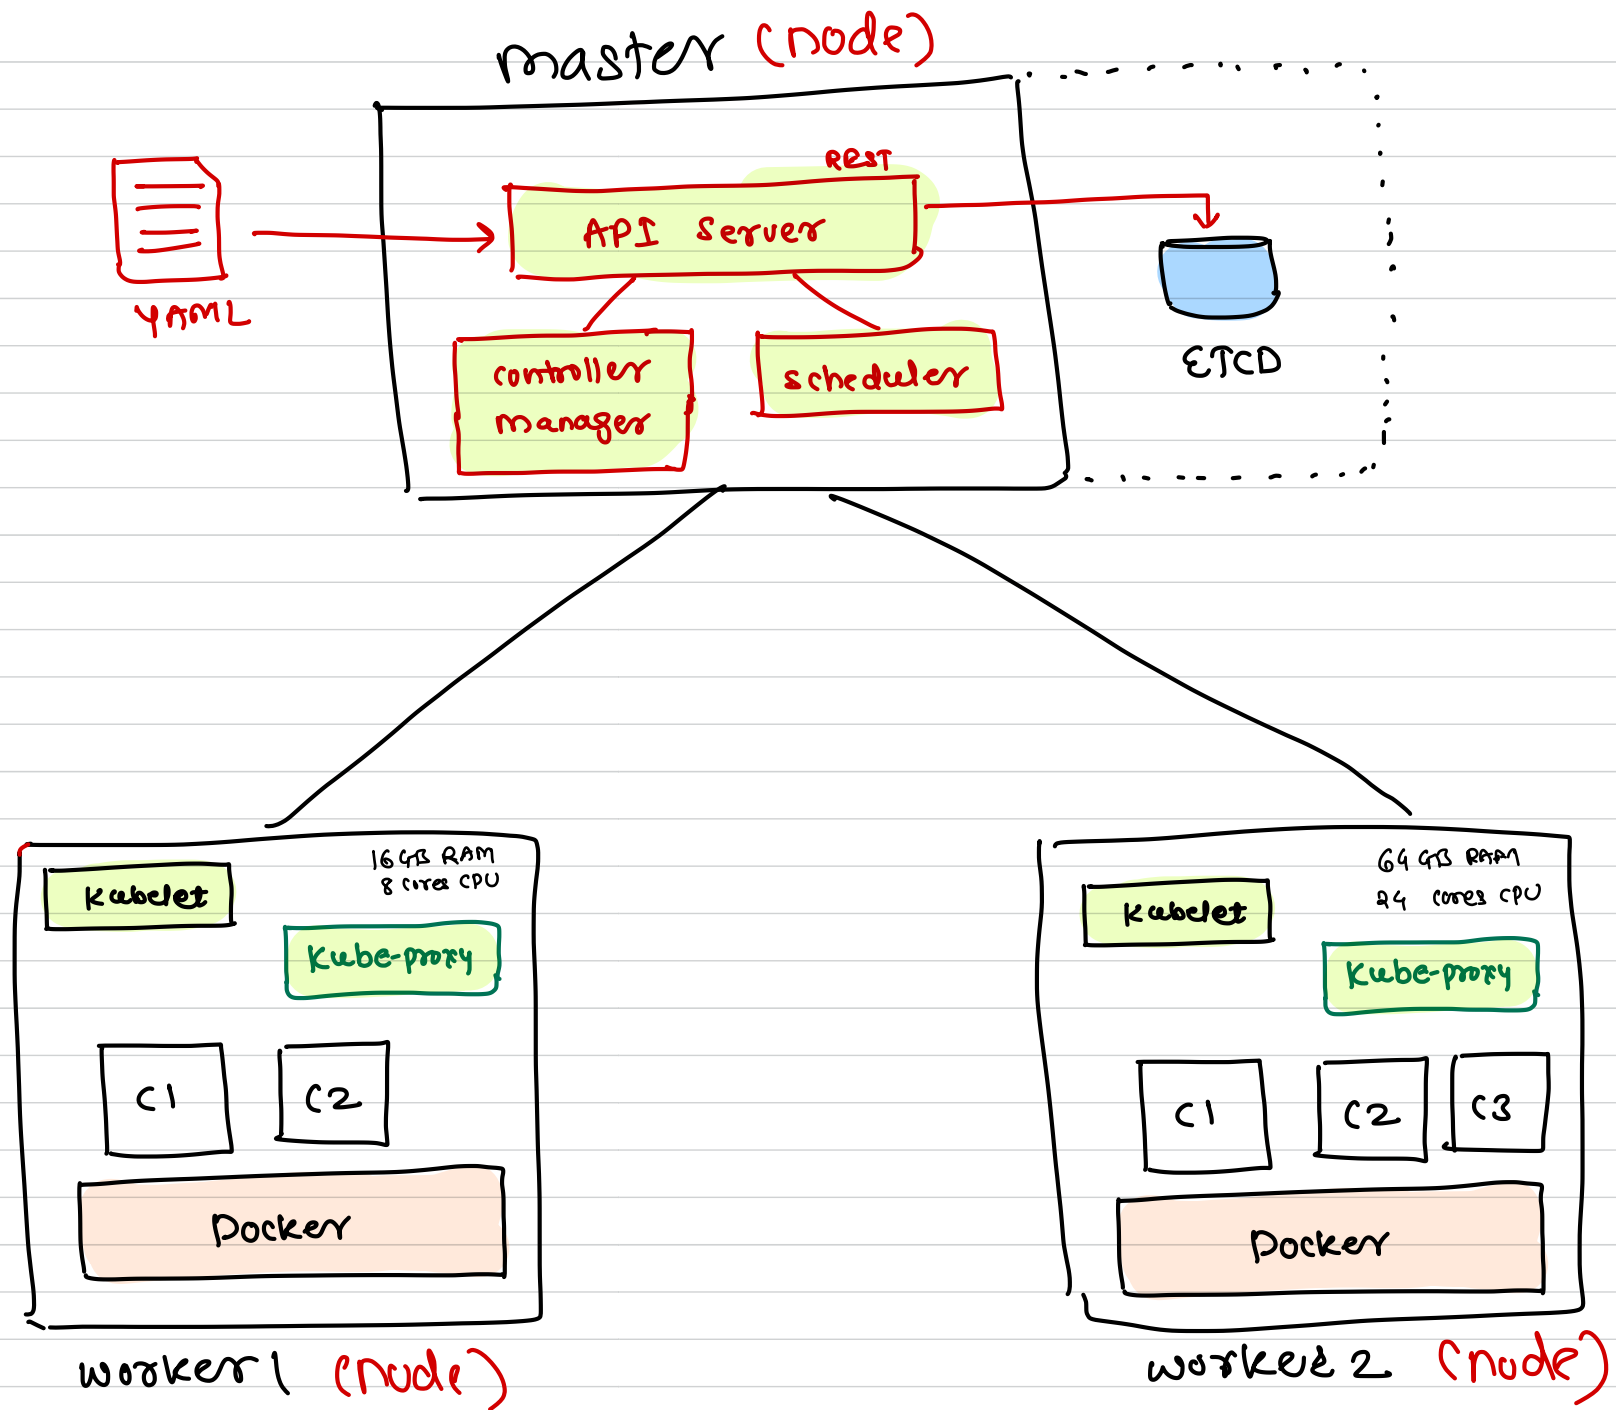
- Provisioning and deployment of containers
- Self healing
- Service discovery and load balancing
- Storage orchestration
- Auto scaling
- Run anywhere self-managed on prem OR managed on cloud
- Automated rollouts and rollbacks
- Secrets and configuration management
- Scale (RC / RS)
- REST APIs at its core
- Security



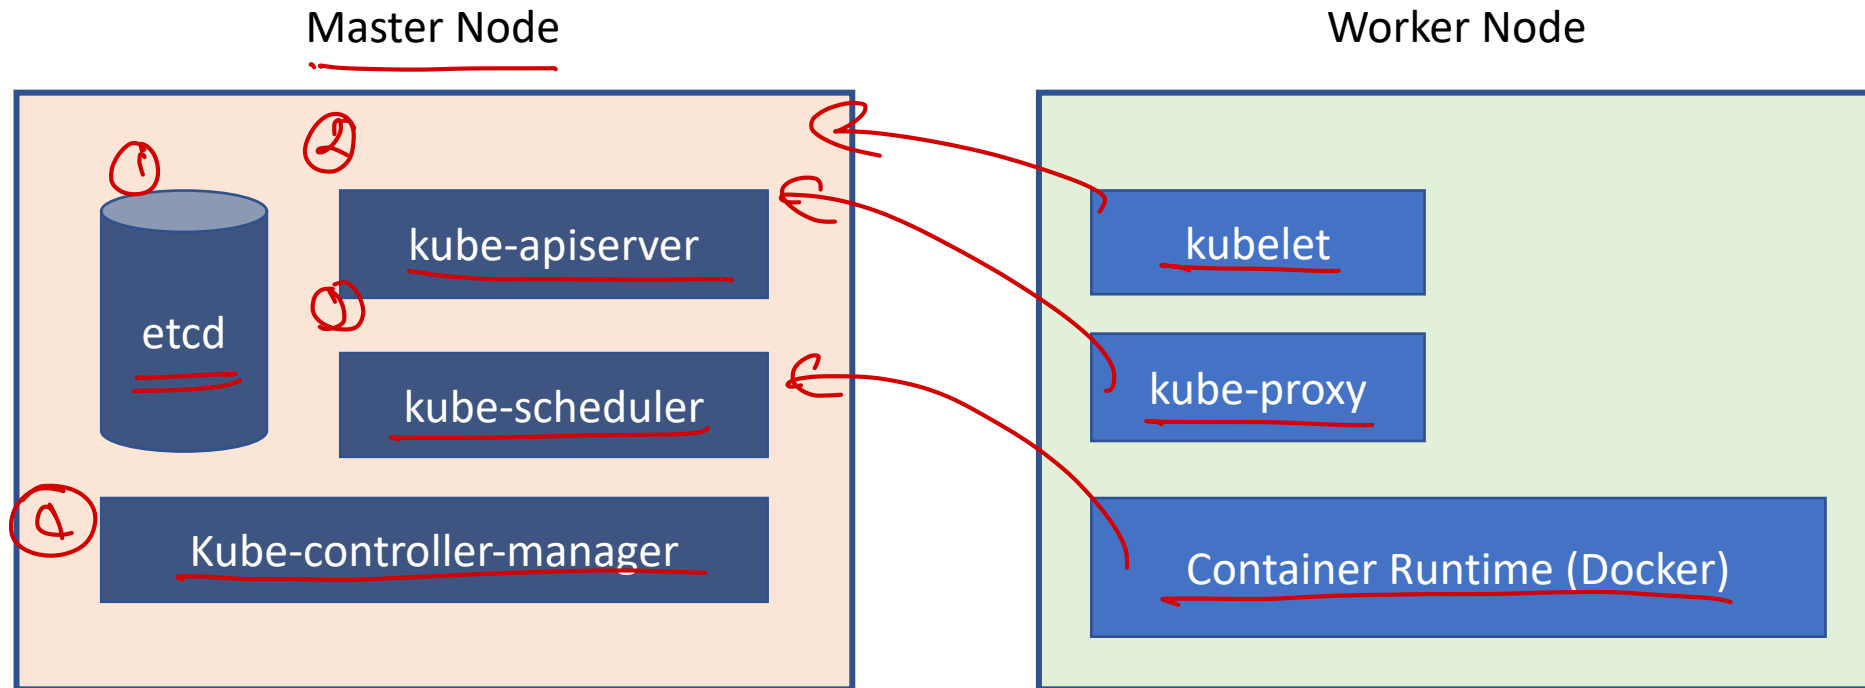
Kubernetes Cluster

- When you deploy Kubernetes, you get a cluster.
- A cluster is a set of machines (nodes), that run containerized applications managed by Kubernetes
- A cluster has at least one worker node and at least one master node
- The worker node(s) host the pods that are the components of the application
- The master node(s) manages the worker nodes and the pods in the cluster
- Multiple master nodes are used to provide a cluster with failover and high availability





Kubernetes Components



Master Components

- Master components make global decisions about the ~~and they~~ detect and respond to cluster events
- Master components can be run on any machine in the cluster
- **kube-apiserver**
 - The API server is a component that exposes the Kubernetes REST API
 - The API server is the front end for the Kubernetes
- **etcd**
 - Consistent and highly-available key value store used as Kubernetes' backing store for all cluster data
 - The configuration data the etcd stores represents the cluster state
 - e.g. which nodes are running the containers, which applications are running etc
- **kube-scheduler**
 - Component on the master that watches newly created pods that have no node assigned, and selects a node for them to run on



Master Components

▪ kube-controller-manager

- Component on the master that runs controllers
- Logically, each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process
- Types
 - **Node Controller**: Responsible for noticing and responding when nodes go down.
 - **Replication Controller**: Responsible for maintaining the correct number of pods for every replication controller object in the system
 - **Endpoints Controller**: Populates the Endpoints object (that is, joins Services & Pods)
 - **Service Account & Token Controllers**: Create default accounts and API access tokens for new namespaces

▪ cloud-controller-manager

- Runs controllers that interact with the underlying cloud providers
- The cloud-controller-manager binary is an alpha feature introduced in Kubernetes release 1.6



Node Components

- Node components run on every node, maintaining running pods and providing the Kubernetes runtime environment
- **kubelet**
 - An agent that runs on each node in the cluster
 - It makes sure that containers are running in a pod
- **kube-proxy**
 - Network proxy that runs on each node in your cluster, implementing part of the Kubernetes service concept
 - kube-proxy maintains network rules on nodes
 - These network rules allow network communication to your Pods from network sessions inside or outside of your cluster
- **Container Runtime**
 - The container runtime is the software that is responsible for running containers
 - Kubernetes supports several container runtimes: Docker, containerd, rktlet, cri-o etc.



Designing a cluster

- Application components and the number of replicas required to run the application smoothly
- Kubernetes v1.16 has following restrictions
 - 5000 nodes
 - 150000 pods
 - 300000 containers
 - 100 pods per node
- Resource requirements of the pods
 - Processor, Memory and HDD requirements for the pods
 - Application Data Storage
- Co-located or external etcd
- Cluster monitoring
- Infrastructure: on-premises, bare metal or managed cloud cluster
- Security



Designing a cluster

- Worker nodes requirements
- Number of workers required to run the pods
- Configuration of worker nodes
- <https://kubernetes.io/docs/setup/best-practices/cluster-large/>

Nodes	Configuration
<u>1-5 nodes</u>	<u>2 vCPU / 4 GB RAM</u>
6-10 nodes	2 vCPU / 7.5 GB RAM
11-100 nodes	4 vCPU / 15 GB RAM
101-250 nodes	8 vCPU / 30 GB RAM
251-500 nodes <i>500</i>	16 vCPU / 30 GB RAM
<u>More than nodes</u>	<u>36 vCPU / 60 GB RAM</u>



Cluster node requirements

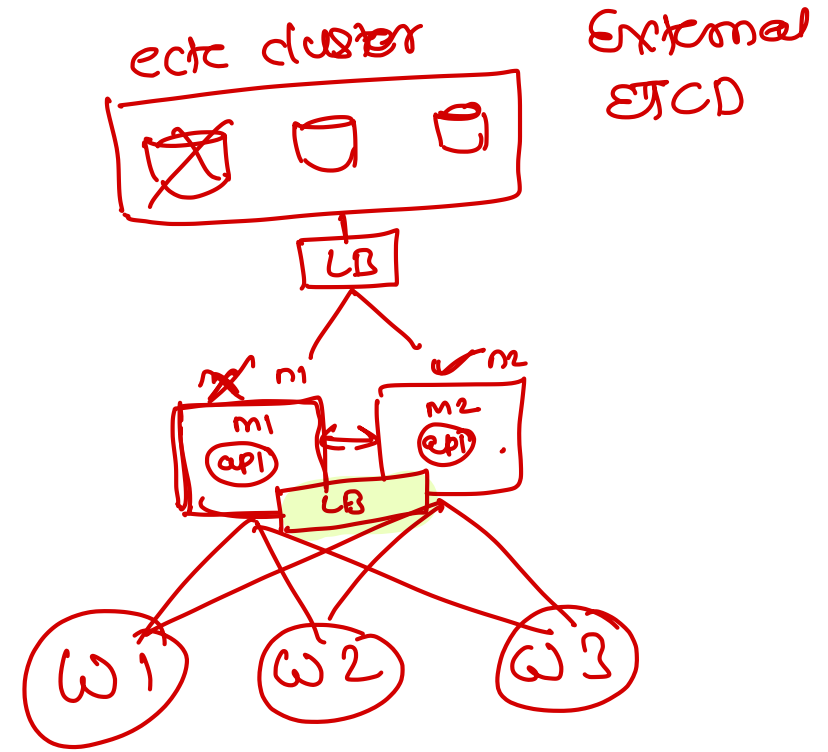
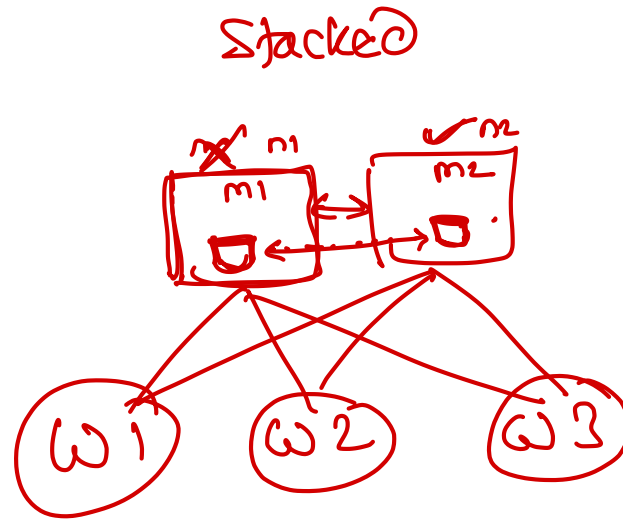
- Cluster can be installed on physical or virtual machine
- Following operating systems are supported with v1.16
 - Ubuntu 16.04+ 18.04
 - Debian 9+
 - Centos 7+
 - RHEL 7+
 - Fedora 25+
 - Container Linux
 - HypriotOS 1.0.1+ + CoreOS
- Minimum configuration on each node is 2CPU and 2 GB RAM
- Control plane ports needed to open
 - 6443, 2379-2380, 10250, 10251, 10252
- Worker nodes ports needed to open
 - 10250, 30000-32767



High Availability Cluster

- No single point of failure
- Kube-apiserver is exposed to worker nodes using a load balancer
- Stacked etcd or External etcd

Raft



Create Cluster

- Use following commands on both master and worker nodes

> sudo apt-get update && sudo apt-get install -y apt-transport-https curl

> curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -

> cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list deb https://apt.kubernetes.io/kubernetes-xenial main EOF

> sudo apt-get update

> sudo apt-get install -y kubelet kubeadm kubect

> sudo apt-mark hold kubelet kubeadm kubect



Initialize Cluster Master Node

- Execute following commands on master node

```
> kubeadm init --apiserver-advertise-address=<ip-address> --pod-network-cidr=10.244.0.0/16
```

```
> mkdir -p $HOME/.kube
```

```
> sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
> sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

- Install pod network add-on

```
> kubectl apply -f
```

```
https://raw.githubusercontent.com/coreos/flannel/2140ac876ef134e0ed5af15c65e414cf26827915/Documentation/kube-flannel.yml
```



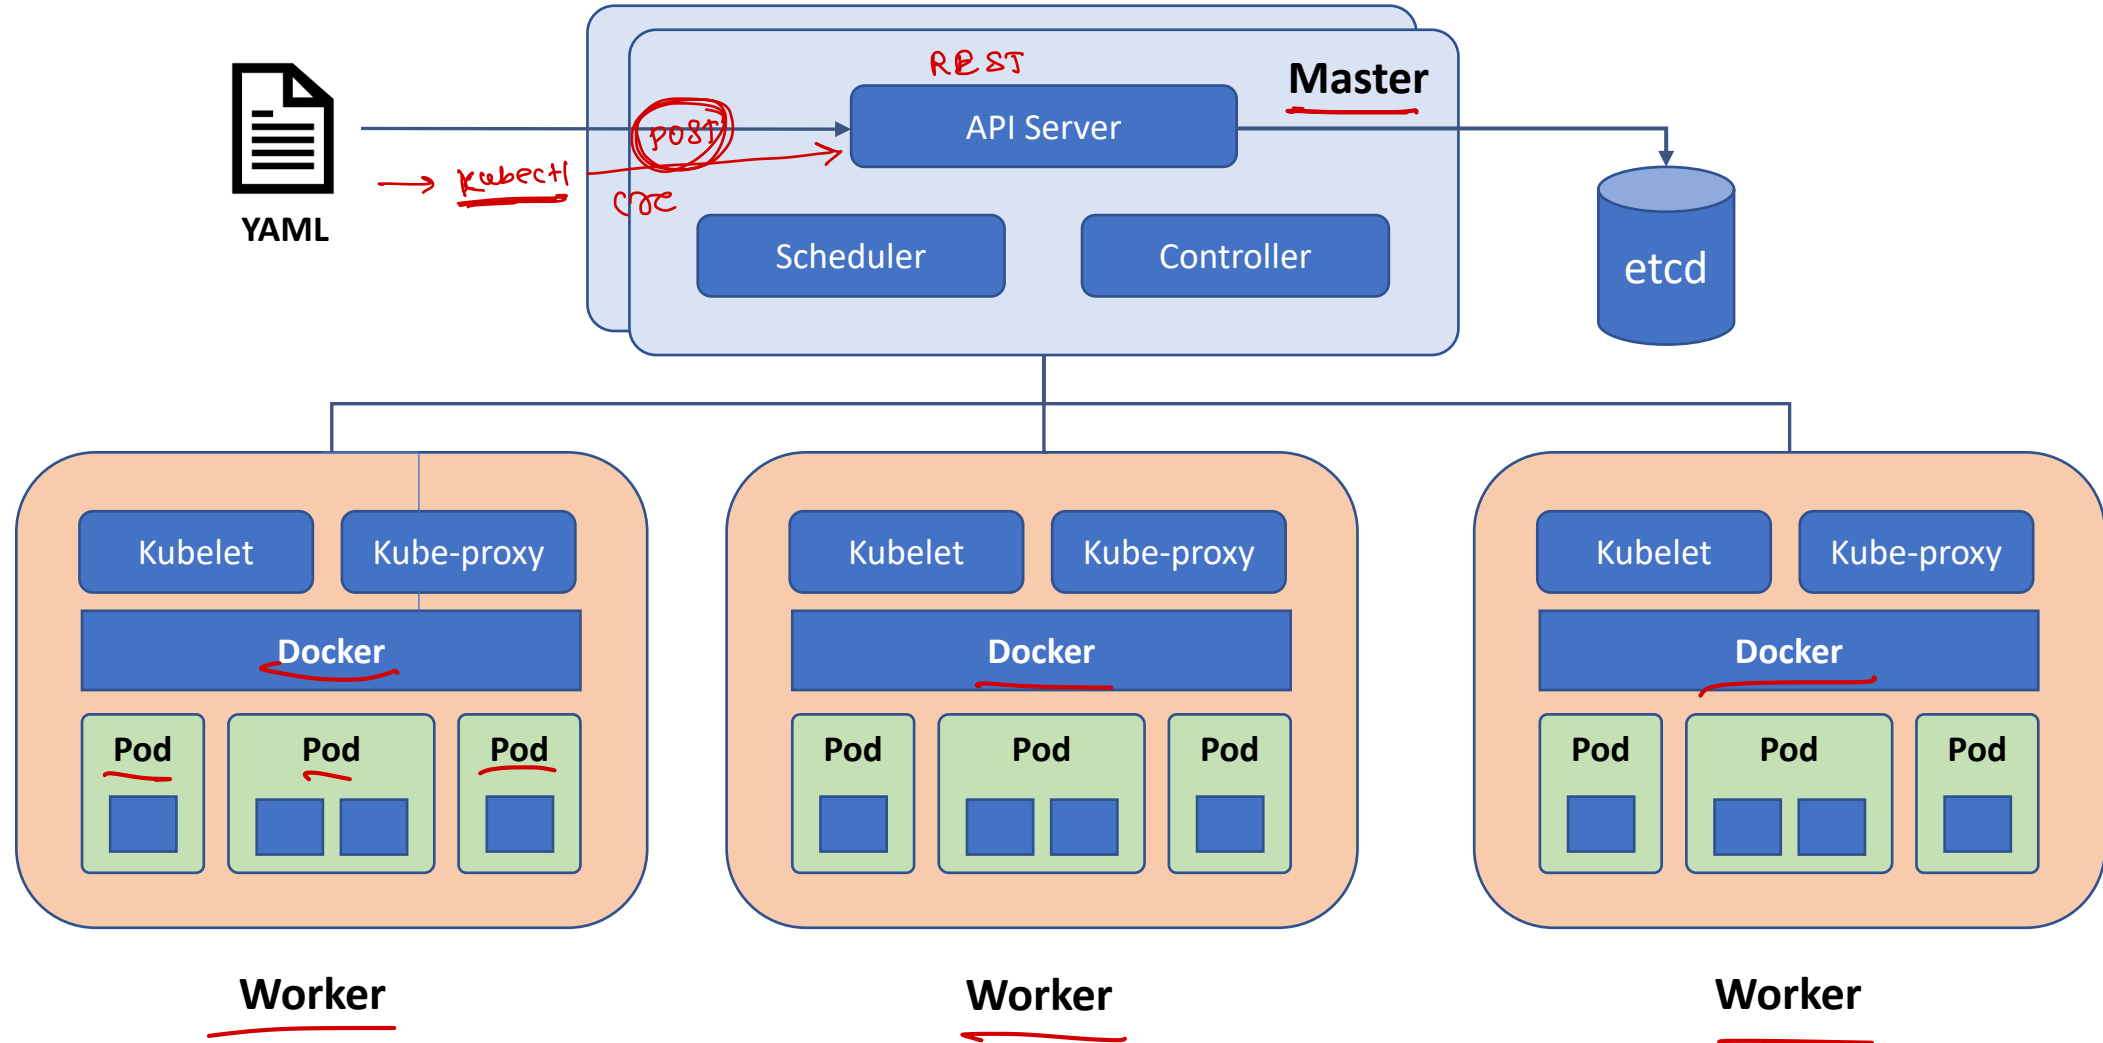
Add worker nodes

- Execute following command on every worker node

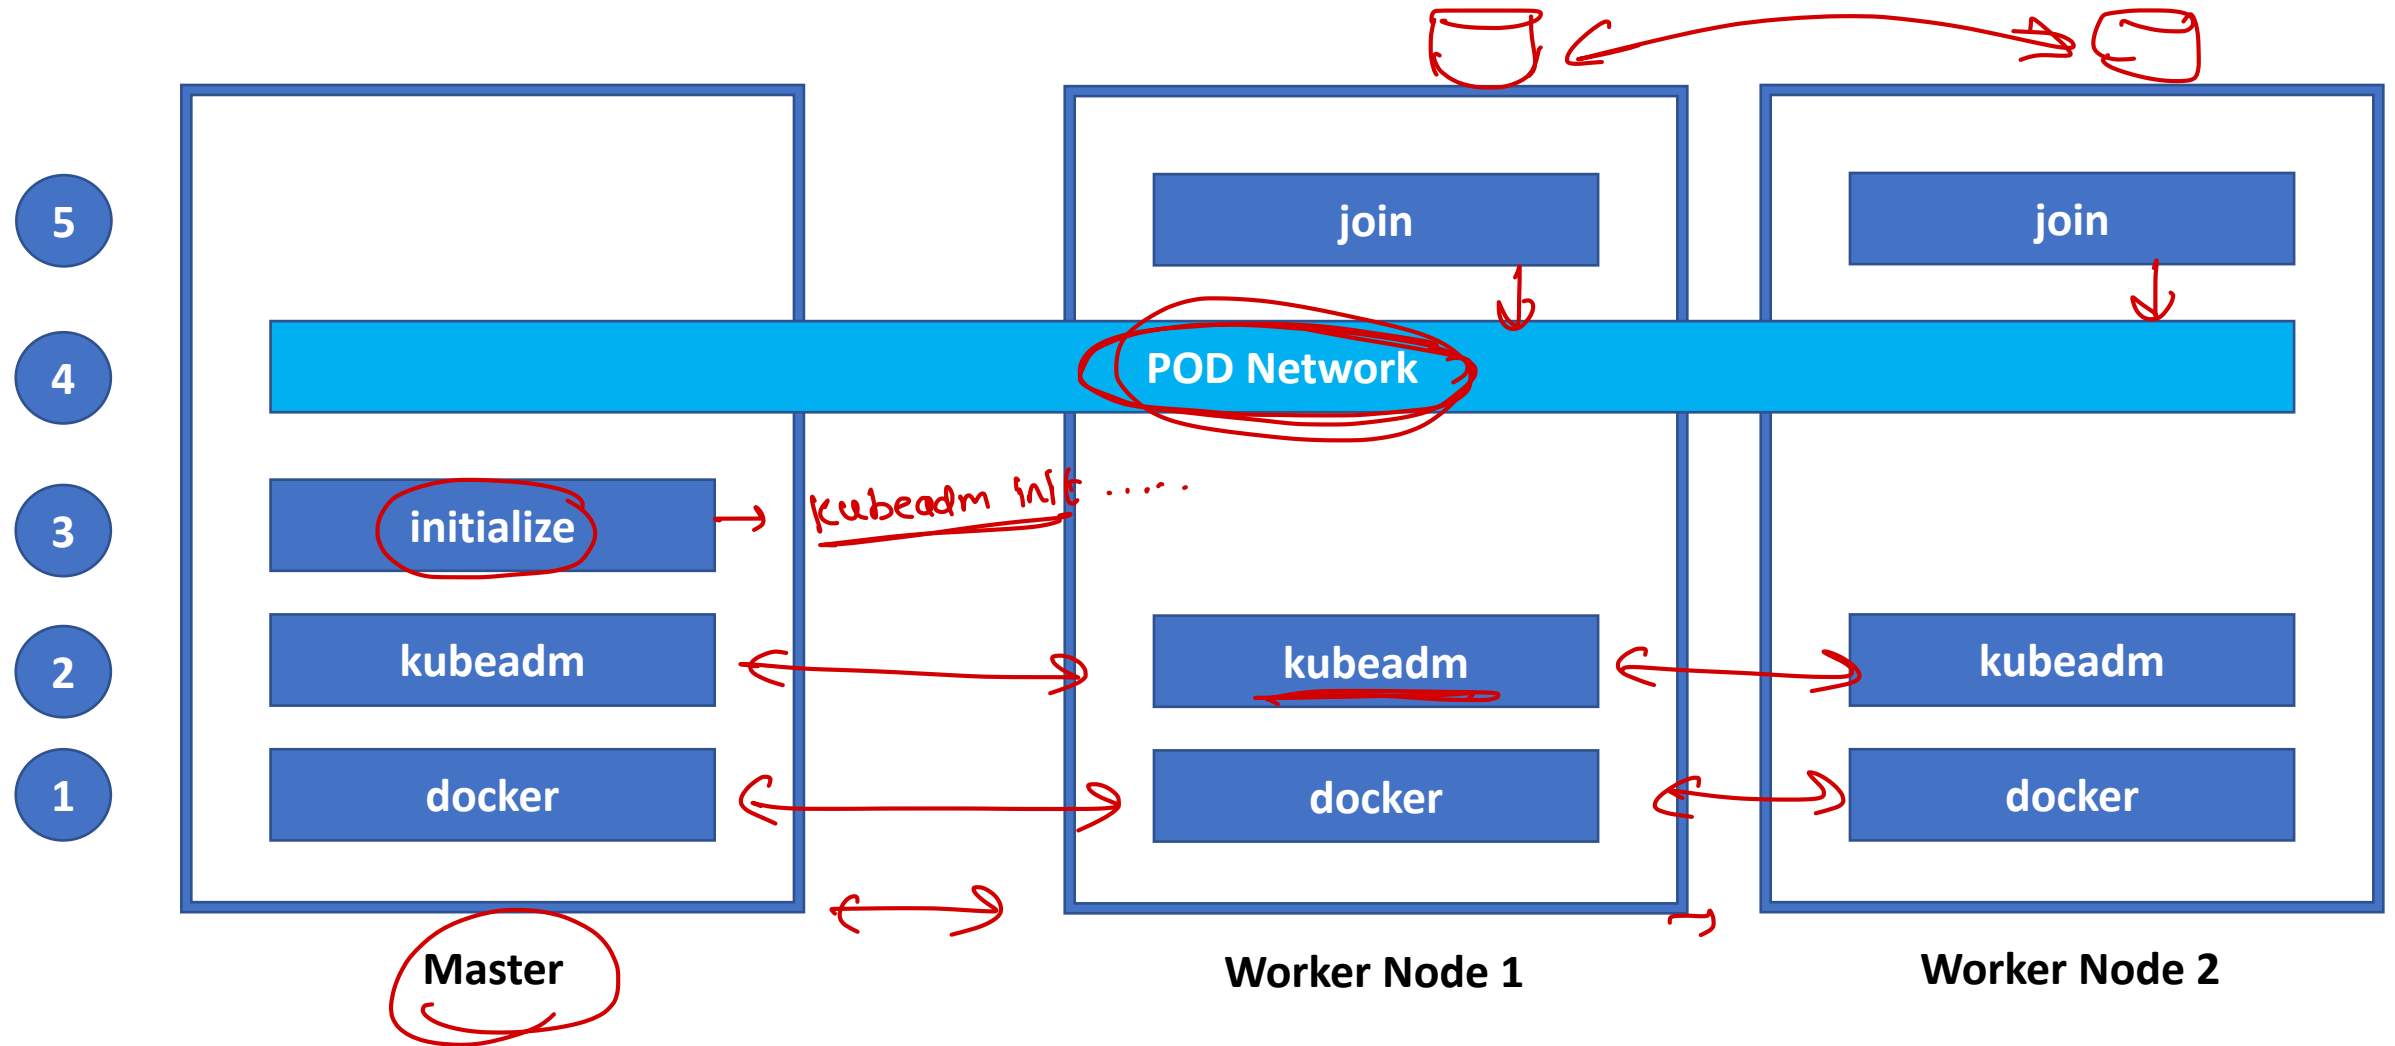
```
> kubeadm join --token <token> <control-plane-host>:<control-plane-port> --discovery-token-ca-cert-hash sha256:<hash>
```



Kubernetes Architecture



Steps to install Kubernetes



Nodes

- Node is not created by Kubernetes. It is created separately and then added to the cluster
- When Kubernetes creates a node, it creates an object to represent the node
- Node validation is done by checking the health of the node
- Valid nodes are eligible for running pods, otherwise they will be ignored while any scheduling activity
- Kubernetes keeps the node in invalid status and keeps checking to see if it becomes valid
- If the node is not required anymore to be the part of cluster, you need to delete it explicitly
- Tasks
 - Get the current cluster info
 - Get the list of nodes in the cluster
 - Find the information about any node



Core Concepts (13%)



Contents

- Understand Kubernetes API primitives
- Create and configure basic Pods



API Model

- Kubernetes API is resource-based (RESTful) programmatic interface provided via HTTP
 - REST APIs are the fundamental fabric of kubernetes
 - All operations, communications between components and commands are handled by the APIs
 - Standard HTTP methods (POST, PUT, PATCH, DELETE, GET)
 - Change notifications on resources via watches
 - Everything in the Kubernetes platform is treated as an API object
 - Kubernetes provides a command like tool (kubectl) which consumes the APIs
 - You can directly call the APIs programmatically
-
- Tasks
 - Get the list of APIs



Resource categories

- Workloads
 - Manage and run the application containers on the cluster
 - Common controllers: Deployments, StatefulSets and Jobs
- Discovery and load-balancer
 - Externally accessible and load balanced services
 - E.g. Services and Ingress
- Config and storage
 - Inject data into the application and keep the data persistent
 - E.g. ConfigMaps, Secrets and Volumes
- Cluster Resources
 - Configuration of the cluster
 - E.g. Cluster Roles, Role Binding and Certificate Signing
- Metadata resources
 - Configuring behavior of the resources within the cluster
 - E.g. HorizontalPodAutoscaler (HPA), PodDisruptionBudget and Notification events



API Terminology

- Resource Type: Name used in the URL
- Resource Kind: Representation of the resource type
- Collection: Instance list of a resource type
- Resource: Single instance of a resource type



API Versioning

- Version is set at the API level rather than at the resource or field level
- Alpha
 - Version name contains alpha e.g. v1alpha1
 - Recommended for use only in short-lived testing clusters
- Beta
 - Version name contains beta e.g. v2beta2
 - Recommended for only non-business-critical uses
- Stable
 - Stable versions e.g. v1, v2, v3 etc
- Tasks
 - Get the list of api versions



API Groups

- Core Group
 - Also known as legacy group
 - apiVersion is used as it is e.g. v1
- Named group
 - apiVersion as \$GROUP_NAME/\$VERSION
 - E.g. batch/v1
- Tasks
 - Get the API groups



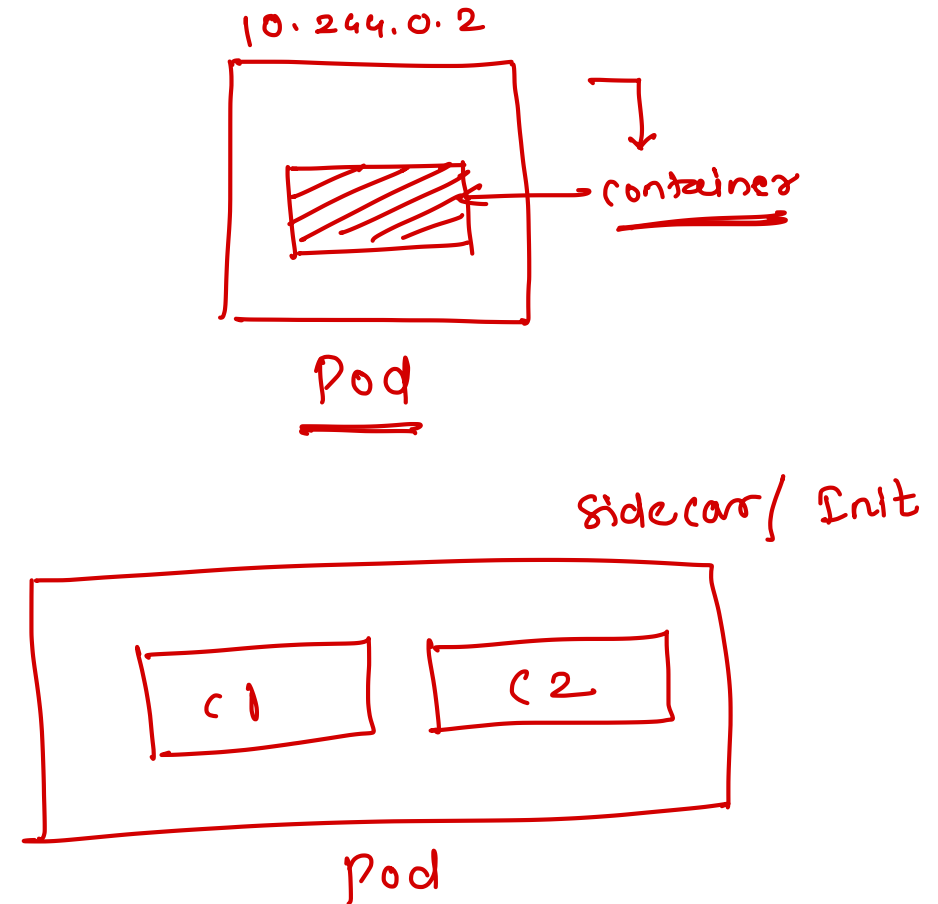
Kubernetes objects

- Persistent entities expressed in YAML format
- Represent the desired state of the cluster
- Describes the following
 - Containerized application running in the cluster
 - Resources made available to these applications
 - Policies for these applications
- Objects can be created using command line utility (kubectl)
- You can also use client libraries to make the API calls directly

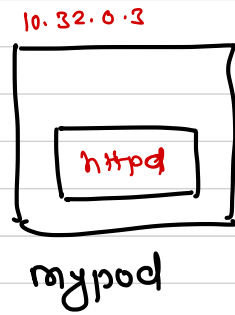


Pods

- Smallest deployable unit in the Kubernetes
- A pod encapsulates
 - Single or multiple containers (Init or App containers)
 - Storage resources
 - Unique Network IP
 - Options governing how the containers should run
 - Ephemeral and disposable entities
- Task
 - Create your first pod using yaml file



```
> kubectl run --image http mypod
```



YAML



YAML

Kubectl

① create

: kubectl create -f <file>

② get

: kubectl get <resource type>

③ describe

: kubectl describe <resource> <name>

④ delete


: kubectl delete <resource> <name>

Phases of Pod


- Phase of a Pod is a simple high level summary showing what state the Pod is in
- Pending
 - Accepted by the cluster
 - Waiting to be scheduled
- Running
 - Scheduled on one of the worker nodes
 - In running/starting/restarting state
- Succeeded
 - All containers in the Pod are terminated successfully and will not be restarted
- Failed
 - At least one container terminated with failure state
- Unknown
 - For some reason the state is not known (may be due to the node communication error)



Container States



- Once the Pod is assigned to a node by scheduler, kubelet starts creating containers using container runtime (like Docker)
- There are three states
 - Waiting ←
 - Running ←
 - Terminated ←



Terminating Pod

- By default, all deletes are graceful within 30 seconds
- kubectl delete commands supports the --grace-period=<seconds> option
- Value of grace-period zero will forcefully delete the pod



Kubernetes Object Management

- The kubectl supports 3 ways to manage the objects
- Imperative commands
 - Operates directly on the live objects
 - Simple with single step commands
 - Do not provide template for creating new objects
 - E.g. `kubectl run nginx --image nginx`
- Imperative object configurations
 - Specify the file containing the full definitions of objects
 - Configuration is simple to understand and can be source controlled
 - Requires additional step of creating YAML file
 - E.g. `kubectl create -f config1.yml`
- Declarative object configuration
 - User does not define the operations to be taken
 - Create, update and delete operations are automatically detected per-object by kubectl
 - Harder to debug and understand results
 - E.g. `kubectl apply -f <directory>/`



Kubernetes Objects - Names

- Every Kubernetes resource requires a unique name within a namespace
- Two different namespaces can have a pod with same name
- Resources have UID as well, which is unique across the whole cluster
- Resource can have names up to 253 characters long
- Allowed characters: digits (0-9), lower case letters (a-z) and dot (.)
- Tasks
 - Find the name of running pod
 - Find uid of a running pod



Namespaces

- Intended for use in environments with many users spread across multiple teams, or projects
- Namespaces provide a scope for resource names
- Names of resources need to be unique within a namespace, but not across namespaces
- Namespaces can not be nested inside one another and each Kubernetes resource can only be in one namespace
- Namespaces are a way to divide cluster resources between multiple users
- When a namespace is deleted, all the resources within the namespace will also get deleted
- Default namespaces
 - default: for objects with no specified namespace
 - kube-public: reserved for cluster usage
 - kube-node-lease: contains lease objects
 - kube-system: objects created by Kubernetes system



Namespaces

- Tasks
 - List all the namespaces
 - Create a new namespace using command
 - Create a new namespace using yaml file



Labels

- Formulate system objects in a structured manner
- They are optionally present in the configuration
- If present
 - They are key-value pairs that are attached to the API objects
 - Each key must be unique in the key-value pairs list
- Tasks
 - Add a pod with labels attached to it

